



Congestion Mitigation (Queue Depth)



Resource Efficiency (KV Cache)



Scaling & Energy Consumption

